


# Operating Systems and Concurrency

## Lab 4 – Graphs and Observations

### Part 2

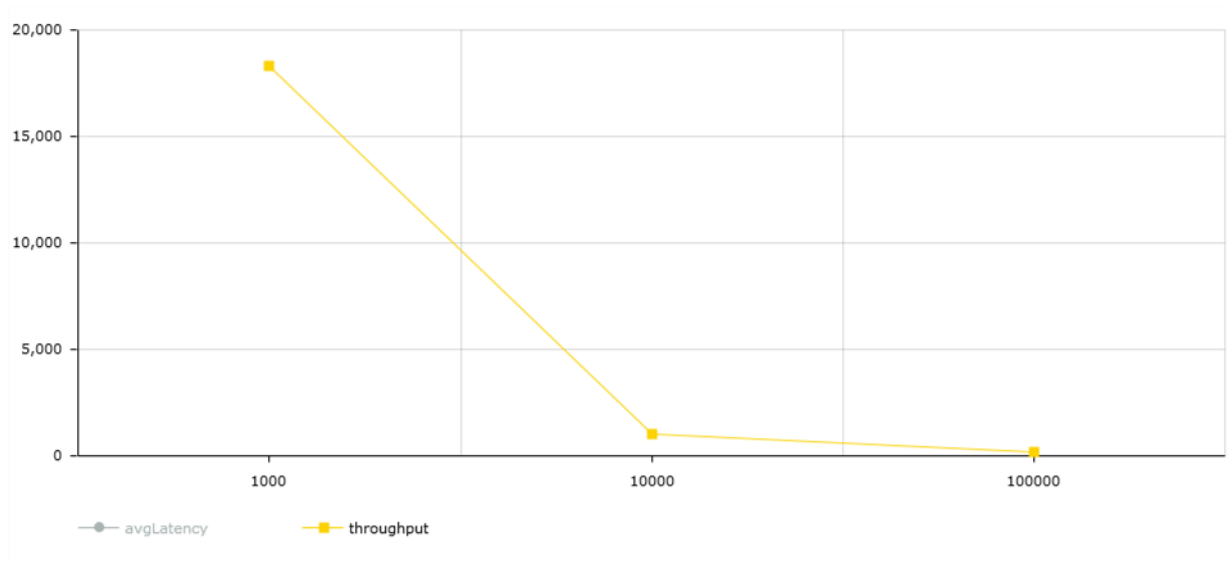
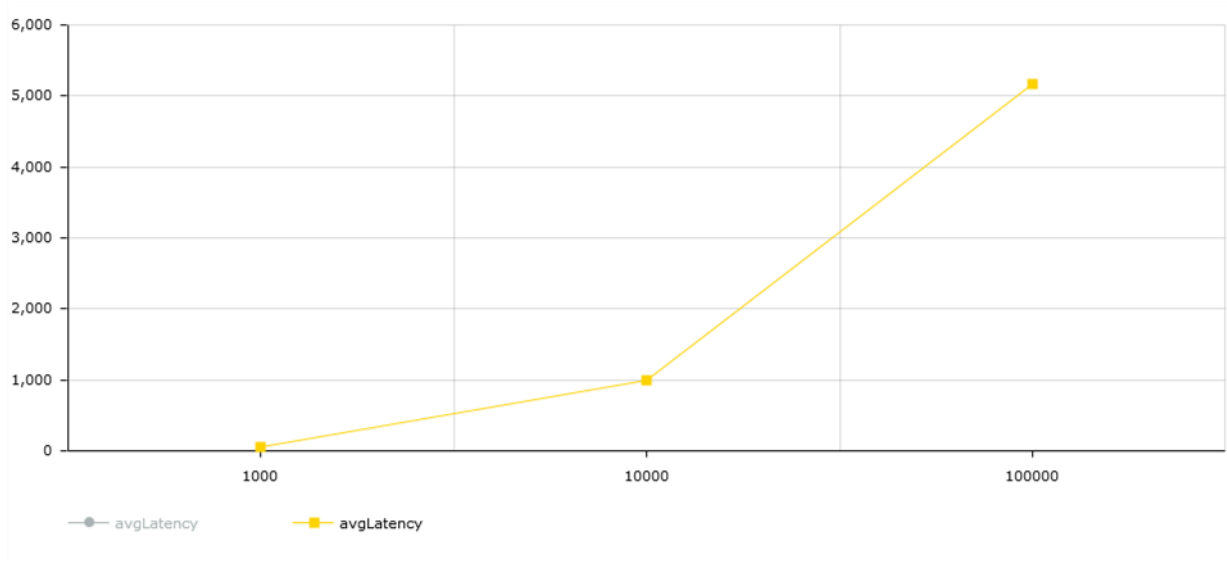
#### 1. Average Latency and Throughput based on N\_DATA:

Throughput decreases and average latency increases, the more data we have.




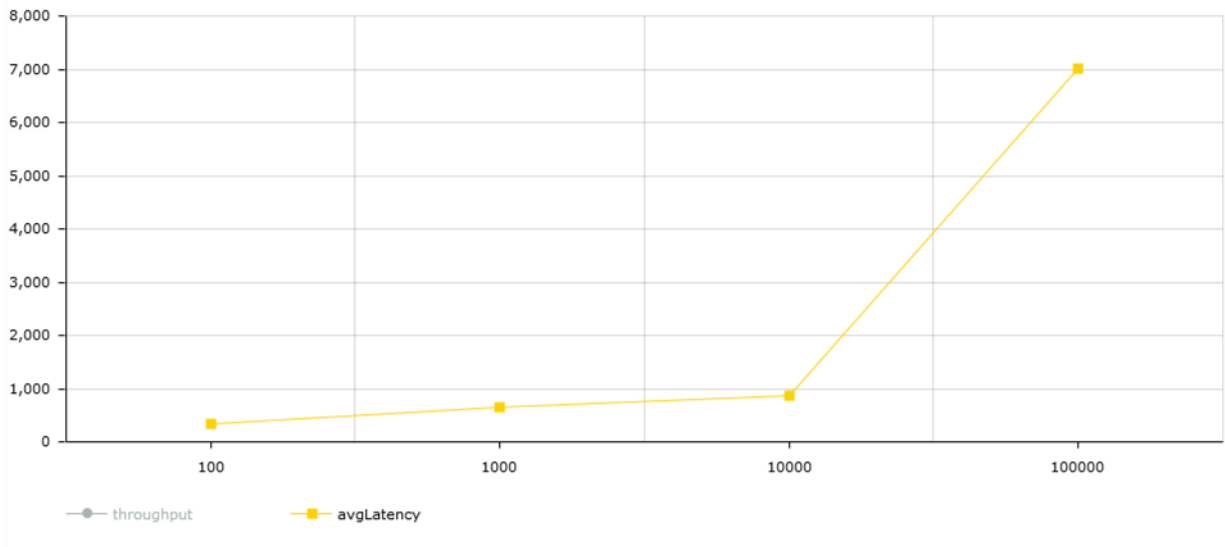
## 2.Average Latency and Throughput depending on the workload of all threads:

Latency increases and throughput decreases the bigger the workload for the threads is.



### 3.Average Latency and Throughput when the first thread has a fixed workload (100.000) and the other 2 threads change workload:

As with the previous test, latency rises and throughput decreases when workload is increased, but the rate at which this happens is lower than before.




# PART 3


## 1.Average Latency and Throughput based on N\_DATA:

Throughput seems to be dropping, while the average latency is going higher..

Throughput chart




Average Latency chart




**2.Average Latency and Throughput depending on the workload of all threads:**

*Throughput chart*




*Average Latency Chart*




**3.Average Latency and Throughput when the first thread has a fixed workload (100.000) and the other 2 threads change workload:**

*Throughput chart*



*Average Latency chart*




# PART 4


## 1.Average Latency and Throughput based on N\_DATA:

Both latency and throughput seem to be pretty constant in this case. There are some minor increments at times, but I would not register those as dramatic and significant.

Average Latency chart




Throughput Chart




**2.Average Latency and Throughput depending on the workload of all threads:**

*Average Latency Chart*




*Thread chart*






### 3.Average Latency and Throughput when the first thread has a fixed workload (100.000) and the other 2 threads change workload:

*Average Latency Chart*



*Throughput Chart*



I believe the program can be optimised further, by separating the input and output buffers into three areas – one for each thread. We will not need mutexes, for this implementation, because the threads will be operating on their own areas only.